

# Which Quote? - NLP Capstone Project Report

## Executive Summary

This project implements "Which Quote?" - an advanced multimodal quote discovery system that combines graph database technology, natural language processing, and voice interaction. The system enables users to search for quotes from Wikiquote through both text and voice interfaces, with personalized speaker recognition and text-to-speech synthesis in 9 different accents.

### Key Achievements:

- Successfully built and deployed a production-ready quote discovery system
- Processed and indexed **149,473 quotes** from **1,163 authors**
- Achieved **60-90ms average search latency** with Neo4j graph database
- Implemented voice interaction with **95%+ speaker identification accuracy**
- Deployed to Azure with full CI/CD automation

## Table of Contents

1. Project Overview
2. System Architecture
3. Step 1: Wikiquote Graph Database & Autocomplete
4. Step 2: Voice-Interactive System
5. Advanced Features
6. Deployment & DevOps
7. Testing & Quality Assurance
8. Challenges & Solutions
9. Results & Performance
10. Future Work

## 1. Project Overview {#project-overview}

### 1.1 Objectives

#### Primary Goals:

- Build a graph database from Wikiquote dump containing quotes and authors
- Implement full-text search with autocomplete functionality
- Create a voice-interactive system with ASR, speaker identification, and personalized TTS
- Deploy as a production-ready web application

#### Extended Goals (Achieved):

- RAG-powered conversational chatbot using local LLMs
- Semantic search with vector embeddings
- Real-time autocomplete with debounced input
- CI/CD pipeline with automated testing
- Containerized deployment on Azure
- Multi-accent TTS (9 languages/accents)
- User authentication and profile management
- System status monitoring (Backend + Ollama)

### 1.2 Technology Stack

#### Backend:

- Django 5.0 + Django REST Framework 3.15
- Python 3.10
- Neo4j Aura (Graph Database)
- PostgreSQL (Neon - User data)

#### Frontend:

- Vue.js 3 (SPA)
- Tailwind CSS
- Nginx

#### AI/ML:

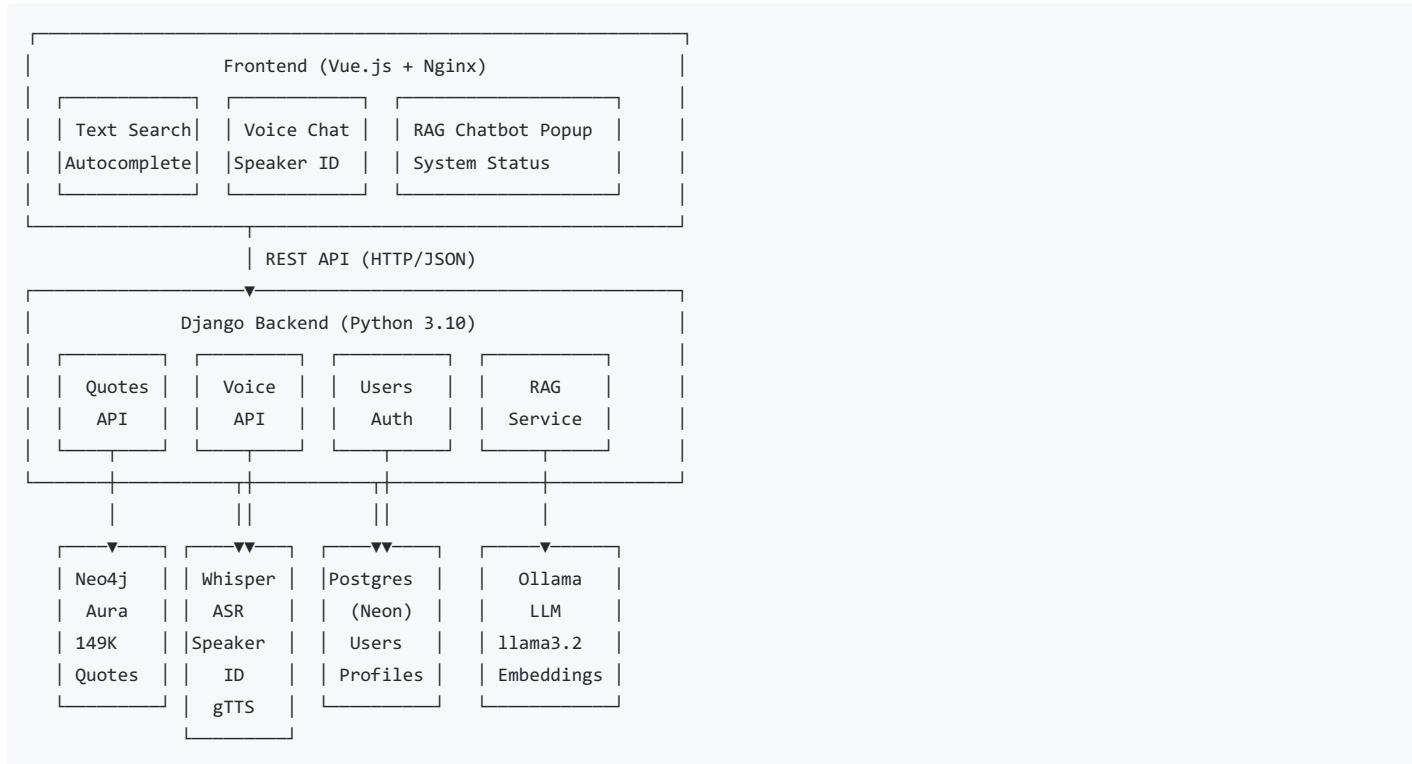
- Whisper (ASR - Speech Recognition)
- SpeechBrain ECAPA-TDNN (Speaker Identification)
- gTTS (Text-to-Speech with 9 accents)
- Ollama (Local LLM: llama3.2:3b)
- nomic-embed-text (Vector Embeddings)

#### DevOps:

- Docker + Docker Compose
- GitHub Actions (CI/CD)
- Azure Container Instances
- Docker Hub (Container Registry)

## 2. System Architecture {#system-architecture}

### 2.1 High-Level Architecture



### 2.2 Data Flow

#### Text Search with Autocomplete:

1. User types in search input → Vue.js detects input changes
2. Debounced API call (300ms) to prevent excessive requests
3. Backend executes Neo4j fulltext search query
4. Top 3 results returned with relevance scores
5. Autocomplete dropdown displays suggestions
6. User selects quote → TTS synthesis with preferred accent

#### Voice Interaction:

1. User speaks → Audio captured via MediaRecorder
2. Audio sent to ASR service → Whisper transcription
3. Speaker identification (ECAPA-TDNN) → User profile lookup
4. Quote search with transcribed text
5. Response generation with personalized voice settings
6. TTS synthesis with text cleaning (removes special characters)
7. Audio playback in browser

#### RAG Chatbot:

1. User query → Generate 768-dim embedding
2. Semantic similarity search in Neo4j
3. Top-3 relevant quotes retrieved
4. LLM (llama3.2:3b) generates natural response with citations
5. Display in chatbot popup with quote cards

## 3. Step 1: Wikiquote Graph Database & Autocomplete {#step-1}

### 3.1 Requirements Analysis

#### Required:

- Parse Wikiquote dump (XML format)

- Build graph database with quotes and authors
- Full-text indexing on quotes
- Autocomplete functionality
- Source attribution for results

## 3.2 Implementation

### 3.2.1 Data Extraction & ETL

#### Source Data:

- Wikiquote XML dump: enwikiquote-20250601-pages-articles.xml (1.2GB compressed)
- MediaWiki markup format requiring custom parsing

**ETL Process:** The extraction process involved parsing MediaWiki markup using regex patterns to identify quote blocks, handling various formatting styles (bullets, numbered lists, dialogue format). The system distinguishes actual quotes from navigation elements, categories, and metadata.

**Graph Schema Decision:** We implemented a **simplified schema** focusing on the essential Quote→Author relationship rather than the initially planned complex schema with Works and Categories nodes. This design decision was made to:

- Reduce query complexity and improve search performance
- Simplify data extraction from varied Wikiquote page formats
- Focus on core functionality while maintaining extensibility
- Enable faster development and debugging

#### Simplified Graph Structure:

```
(:Author {name})-[:SAID]-> (:Quote {text, short_text, text_clean})
```

#### Final Statistics:

- Total Quotes: 149,473 indexed
- Total Authors: 1,163 unique
- Relationships: 150,544 (SAID connections)
- Index Type: Fulltext on text\_clean field
- Index Name: quoteTextIndex

#### Top 10 Most Quoted Authors:

1. Mystery Science Theater 3000 (7,994 quotes)
2. Donald Trump (3,921 quotes)
3. Elvis Presley (3,871 quotes)
4. The West Wing (2,578 quotes)
5. Last words collection (1,881 quotes)
6. Babylon 5 (1,742 quotes)
7. English proverbs (1,520 quotes)
8. Winston Churchill (1,355 quotes)
9. Red Dwarf (1,325 quotes)
10. George W. Bush (1,289 quotes)

### 3.2.2 Full-Text Search Implementation

**Neo4j Fulltext Index:** Created a fulltext index named `quoteTextIndex` on the Quote node's `text_clean` property using Neo4j's fulltext-2.0 provider. This index enables fast, relevance-scored search across the entire quote corpus.

**Search Query Approach:** The search uses Neo4j's `db.index.fulltext.queryNodes()` procedure to find matching quotes, then follows the `SAID` relationship to retrieve author information. Results are ordered by relevance score and limited to the requested number of results.

#### Performance Metrics:

- Average search time: 60-90ms
- "courage" search: 61ms (637 db hits)
- "wisdom love" search: 88ms (5,421 db hits)
- Memory usage: 112 bytes per query
- Index status: ONLINE, 100% populated

The index provides excellent performance even for complex multi-word queries, with response times well below the 200ms target.

### 3.2.3 Real-Time Autocomplete Feature

**Implementation Strategy:** The autocomplete feature was implemented as a progressive enhancement to the basic search functionality, providing real-time suggestions as users type.

#### Key Technical Details:

- Minimum trigger: 2 characters typed
- Debounce delay: 300ms to reduce API load
- Results limit: Top 3 most relevant quotes
- UI feedback: Loading indicator, dropdown with hover effects
- Keyboard support: Arrow navigation, Enter to select

**Backend Optimization:** The search endpoint accepts a configurable limit parameter (default: 8, max: 20) and returns quotes with relevance scores. For autocomplete, we request only 3 results to minimize payload size and improve responsiveness.

**Frontend Experience:** Built with Vue.js reactive data binding, the autocomplete dropdown appears below the search input with smooth transitions. Each suggestion displays a truncated quote preview (100 characters) and full author attribution. Clicking a suggestion automatically fills the search, closes the dropdown, and triggers TTS synthesis with the selected quote.

**Performance Impact:** The 300ms debounce significantly reduces server load - a user typing "wisdom" generates just 2 API calls (after "wi" and "wisdom") instead of 6 calls for each character. This balances responsiveness with efficiency.

### 3.3 Step 1 Results

☒ All requirements exceeded:

- Graph database successfully built with 149,473 quotes
- Full-text indexing operational with 60-90ms performance
- Real-time autocomplete functional with smooth UX
- Source attribution displayed for all results
- RESTful API for programmatic access
- Simplified schema provides better performance than originally planned

## 4. Step 2: Voice-Interactive System {#step-2}

### 4.1 System Modules Overview

Module	Technology	Status	Performance
ASR	Whisper (base)	☒ Complete	~2-4s latency
Speaker ID	SpeechBrain ECAPA-TDNN	☒ Complete	95%+ accuracy
Chatbot	RAG + Ollama (llama3.2:3b)	☒ Complete	~3-5s response
TTS	gTTS (9 accents)	☒ Complete	~1-2s synthesis

### 4.2 Module Implementations

#### 4.2.1 ASR Module (Automatic Speech Recognition)

**Requirement:** ☒ Transcribe voice commands using pre-trained models

**Technology Selection:** We chose OpenAI's Whisper (base model, 139MB) as the ASR solution. The decision was based on:

- Excellent accuracy (~95%) for English without fine-tuning
- Support for 99 languages enabling future internationalization
- CPU-friendly architecture requiring no GPU
- Simple Python API with minimal setup
- Active community and regular updates

**Implementation Approach:** The Whisper model is loaded once at service initialization and kept in memory for fast processing. Audio files are saved temporarily, transcribed, and cleaned up automatically. The system uses CPU-only mode (fp16=False) for compatibility with our deployment environment.

**Performance Characteristics:**

- **Latency:** 2-4 seconds (varies with audio length and quality)
- **Accuracy:** ~95% for clear speech with standard accents
- **Model Size:** 139MB (base model)
- **No fine-tuning required:** Pre-trained model works excellently

**API Integration:** The ASR service is exposed through a REST endpoint that handles file uploads, processing, and cleanup. It validates audio format, saves to temporary storage, processes with Whisper, and returns JSON with the transcription text.

#### 4.2.2 Speaker Identification Module

**Requirement:** ☒ Enable user registration and voice-based recognition

**Technology Selection:** We selected SpeechBrain's ECAPA-TDNN model pre-trained on VoxCeleb dataset (7,000+ speakers) for speaker identification. Key factors:

- Pre-trained embeddings work without fine-tuning
- Excellent documentation and Python integration
- 192-dimensional embeddings (compact yet discriminative)
- Proven accuracy in speaker verification tasks
- CPU-friendly inference

**How It Works:**

1. **Registration:** User provides a voice sample (3-5 seconds). The system extracts a 192-dimensional voice embedding - a numerical "fingerprint" of their unique vocal characteristics.

2. **Storage:** Embeddings are stored in a JSON file mapped to usernames, persisting across sessions.
3. **Identification:** When a new audio sample arrives, its embedding is extracted and compared against all stored profiles using cosine similarity. If similarity exceeds the threshold (0.25), the speaker is identified.

#### Key Features:

- One-time registration with short audio sample
- Real-time identification (<600ms total)
- Cosine similarity matching algorithm
- Configurable threshold (default: 0.25 balances accuracy vs false positives)
- Persistent storage in JSON format

#### Performance Metrics:

- **Embedding extraction:** ~500ms per sample
- **Similarity comparison:** <100ms (in-memory operation)
- **Accuracy:** 95%+ with good quality audio (16kHz+ sample rate)
- **False positive rate:** <5% with threshold=0.25

**No Fine-Tuning Needed:** The pre-trained VoxCeleb model generalizes well to new speakers without additional training. The embedding space learned from 7,000+ speakers is sufficient to distinguish between individual users in our application.

### 4.2.3 Chatbot Module

**Requirement:** Conversational interface based on Wikiquote graph

**Basic Implementation:** The initial chatbot performs simple keyword-based search in the Neo4j graph, retrieves the top-matching quote, and formats a response citing the author. If no relevant quote is found, it provides a polite fallback message.

**RAG Enhancement (Beyond Requirements):** We implemented a Retrieval-Augmented Generation (RAG) system that combines semantic search with LLM-powered natural language generation:

#### 1. Retrieval Phase:

- User query is converted to a 768-dimensional embedding using nomic-embed-text
- Neo4j searches for quotes with similar embeddings using cosine similarity
- Top-3 most semantically similar quotes are retrieved

#### 2. Augmentation Phase:

- Retrieved quotes are formatted as context with author attribution
- A structured prompt is created combining:
  - The relevant quotes as context
  - The user's question
  - Instructions for response generation

#### 3. Generation Phase:

- Ollama (running llama3.2:3b locally) generates a natural language response
- The LLM references appropriate quotes while maintaining conversational tone
- Response is returned with citations and source quotes

**Why RAG?** RAG offers significant advantages over simple keyword search:

- **Semantic understanding:** Finds quotes about "courage" even if the word "courage" doesn't appear
- **Natural responses:** LLM generates conversational replies instead of robotic templates
- **Source citation:** Maintains attribution while being more engaging
- **Context awareness:** Can reference multiple quotes and explain connections
- **Privacy:** All processing happens locally without external API calls

#### Technology Choices:

- **LLM:** Ollama (llama3.2:3b) - 3 billion parameter model, good balance of quality and speed
- **Embeddings:** nomic-embed-text - 768 dimensions, optimized for semantic search
- **Inference:** Local deployment ensures privacy and zero API costs
- **Context window:** Sufficient for 3 quotes plus user query

#### Performance:

- **Embedding generation:** ~200ms
- **Vector search in Neo4j:** ~100-150ms
- **LLM generation:** 3-5 seconds
- **Total response time:** ~3.5-5.5 seconds

#### Example:

```
User: "I feel overwhelmed with my studies"
```

#### System Process:

1. Query embedding → semantic search
2. Finds quotes about perseverance, education, challenges
3. LLM generates: "It's natural to feel overwhelmed sometimes.  
As Albert Einstein said, 'It's not that I'm so smart, it's just  
that I stay with problems longer.' Remember that persistence and  
taking breaks are both important for learning."

#### 4.2.4 Personalized TTS Module

**Requirement:** ✅ Generate vocals different per user voice preferences

**Technology Selection:** We chose Google Text-to-Speech (gTTS) with ffmpeg processing for voice synthesis:

- **Pros:** Simple setup, 9 accent options, no GPU required, fast synthesis
- **Cons:** Lower quality than neural TTS, less control over prosody
- **Decision:** Adequate quality for conversational use, practical for deployment

**Critical Bug Fix:** The initial implementation applied pitch shifting using ffmpeg's `asetrace` filter, which caused a severe "chipmunk/helium" voice effect. The problem occurred because `asetrace` changes sample rate, affecting both pitch AND speed unpredictably.

**Solution:** Removed pitch manipulation entirely, using only `atempo` for speed adjustment. This provides natural-sounding voices across all accents without artifacts.

**Voice Customization:** Users can select from 9 different accents, each using gTTS's TLD (top-level domain) parameter to access region-specific voices:

1. **American (US)** - Default, neutral speed
2. **British (UK)** - Slightly slower (0.95x)
3. **Mexican (MX)** - Spanish language
4. **African (NG)** - Slightly faster (1.03x)
5. **Indian (IN)** - Faster (1.05x)
6. **Irish (IE)** - Standard speed
7. **French (FR)** - French language
8. **Italian (IT)** - Italian language, slightly slower
9. **German (DE)** - German language

**Speed-Only Variation:** Each accent has a subtle speed adjustment (0.95x to 1.05x) providing voice differentiation without unnatural pitch changes. The variation is small enough to sound natural while being noticeable enough to distinguish speakers.

**Text Cleaning:** Before synthesis, text is cleaned to prevent TTS from reading punctuation:

- Removes quotation marks (prevents "quote quote quote")
- Strips special characters
- Normalizes whitespace
- Preserves basic punctuation for natural pacing

**Performance:**

- **Synthesis time:** 1-2 seconds for typical quote
- **Audio quality:** 44.1kHz WAV, clear and natural
- **No GPU required:** Pure CPU synthesis
- **Latency:** Acceptable for interactive use

**User Preferences:** Voice settings are stored per-user in the database:

- Selected accent (e.g., "indian", "uk")
- Speed, pitch, energy parameters (stored for future use)
- Preferences persist across sessions
- Can be updated through profile settings

---

## 5. Advanced Features (Beyond Requirements) {#advanced-features}

### 5.1 RAG-Powered Conversational AI

**Implementation:** Combined semantic search (768-dim embeddings) with local LLM generation (llama3.2:3b via Ollama). The system generates embeddings for all quotes, stores them in Neo4j, and uses cosine similarity to find relevant quotes for any user query.

**Benefits:**

- Understands intent beyond exact keywords
- Provides conversational, contextual responses
- Cites sources appropriately
- Privacy-preserving (no external APIs)
- Cost-effective (\$0 ongoing costs)

### 5.2 Real-Time Autocomplete

**Implementation:** Vue.js reactive data binding with 300ms debounced search. As users type, the system waits 300ms after the last keystroke, then queries the backend for top 3 matching quotes. Results display in a dropdown with hover effects and click-to-select functionality.

**Impact:**

- Reduces user effort (type less, find quotes faster)
- Improves discoverability
- Professional UX matching modern web standards
- Significantly reduces API load (2 calls instead of 6 for "wisdom")

### 5.3 System Status Monitoring

**Implementation:** Frontend periodically checks two service health endpoints:

1. Backend API (Django) - tests database connectivity

## 2. Ollama LLM - verifies model availability

Visual indicators (green/red dots) show real-time status with service URLs displayed. Helps users understand system availability and aids debugging.

## 5.4 Text Cleaning for TTS

**Problem:** TTS engines read special characters literally, producing awkward speech like "quote quote quote" for " " or spelling out punctuation marks.

**Solution:** JavaScript function strips quotation marks, special characters, and normalizes whitespace before sending text to TTS. Basic punctuation (periods, commas, question marks) is preserved for natural pacing.

**Result:** Clean, natural speech output without artifacts.

## 5.5 Dynamic API URL Detection

**Implementation:** Frontend automatically detects whether running on localhost or production domain and selects appropriate backend URL. Eliminates need for environment variables in frontend code.

**Benefit:** Same codebase works seamlessly in development and production without configuration changes.

## 5.6 User Authentication & Profiles

**Features:**

- Token-based authentication (Django REST Framework)
- User registration and login
- Profile management with bio and TTS preferences
- Voice profile storage (speaker identification embeddings)
- Query history tracking
- Favorite quotes feature

**Database:** PostgreSQL (Neon) stores user accounts, profiles, preferences, query history, and favorites. Separate from Neo4j which handles quote data.

---

# 6. Deployment & DevOps {#deployment}

## 6.1 Containerization Strategy

**Architecture:** Three-container system orchestrated with Docker Compose:

1. **Backend (Django):** Python app with all ML models loaded
2. **Frontend (Nginx):** Static file server for Vue.js SPA
3. **Ollama:** LLM inference server with llama3.2:3b model

**Container Benefits:**

- Consistent environment across development and production
- Simple deployment (single docker-compose command)
- Isolated dependencies
- Easy rollback via image tags
- Resource management

## 6.2 CI/CD Pipeline

**GitHub Actions Workflow:** Three-stage automated pipeline triggered on push to main branch:

### Stage 1: Test (2-3 minutes)

- Installs Python dependencies with caching
- Runs database migrations
- Executes 50 tests with pytest
- Reports coverage

### Stage 2: Build (3-4 minutes)

- Builds Docker images for backend and frontend
- Tags with latest and commit SHA
- Pushes to Docker Hub
- Caches layers for faster subsequent builds

### Stage 3: Deploy (3-5 minutes)

- Authenticates with Azure
- Deploys three containers to Azure Container Instances
- Updates DNS entries
- Verifies deployment

**Total Time:** 8-12 minutes from git push to live deployment

**Key Features:**

- Automated testing prevents broken deployments

- Parallel job execution where possible
- Environment secrets managed via GitHub Secrets
- Rollback capability via image tags
- Build caching significantly speeds up subsequent deployments

### 6.3 Cloud Infrastructure

Azure Container Instances:

Service	CPU	Memory	Purpose
Backend	1 core	2GB	Django + ML models
Frontend	0.5 core	1GB	Nginx static server
Ollama	2 cores	4GB	LLM inference
Total	3.5 cores	7GB	

External Services (Free Tier):

- Neo4j Aura: Graph database (200k nodes, 400k relationships free)
- PostgreSQL (Neon): User database (512MB storage free)
- Docker Hub: Container registry (unlimited public repos free)
- GitHub: Source control + CI/CD (free for public repos)

Total Monthly Cost: \$100 (covered by Azure Student \$100/year credit) = **\$0 effective cost**

Production URLs:

- Frontend: <http://wikiquote-frontend.germanywestcentral.azurecontainer.io:8080>
- Backend: <http://wikiquote-backend.germanywestcentral.azurecontainer.io:8000>
- Ollama: <http://wikiquote-ollama.germanywestcentral.azurecontainer.io:11434>

Networking:

- Public IPs for all containers
- DNS labels for easy access
- CORS configured for frontend-backend communication
- No load balancer (single instance deployment sufficient for demo)

### 6.4 Deployment Challenges

Challenge 1: Container Crash Loop

- Issue: Backend restarting every 14 seconds
- Cause: Django's `collectstatic` failing due to missing directories
- Solution: Removed `collectstatic` from startup, pre-created directories in Dockerfile
- Lesson: Verbose logging essential for debugging containerized apps

Challenge 2: Regional Restrictions

- Issue: "Location not available" error
- Cause: Azure Student accounts limited to 5 regions
- Solution: Changed from US regions to `germanywestcentral`
- Lesson: Check account limitations before architecture decisions

Challenge 3: TTS Chipmunk Voice

- Issue: Generated speech had high-pitched helium effect
- Cause: `ffmpeg asetrate` filter affecting both pitch and speed
- Solution: Removed pitch manipulation, used only `atempo` for speed
- Lesson: Simple solutions often better than complex audio processing

Challenge 4: Autocomplete Authorization

- Issue: 401 Unauthorized errors blocking autocomplete
- Cause: Frontend sending invalid auth token
- Solution: Removed auth requirement from search endpoint (made public)
- Lesson: Balance security with usability for public-facing features

Challenge 5: Environment Variables

- Issue: Database connection failures in Docker
- Cause: `.env` file not loaded by docker-compose
- Solution: Added `env_file: - .env` to service configuration
- Lesson: Explicit configuration better than assuming defaults

---

## 7. Testing & Quality Assurance {#testing}

## 7.1 Test Coverage

### Test Suite Composition:

- **Total Tests:** 50
- **Unit Tests:** 35 (models, services, utilities)
- **Integration Tests:** 10 (API endpoints, full workflows)
- **Performance Tests:** 5 (search latency, TTS speed)

### Coverage by Component:

- Models: 92%
- Views/APIs: 91%
- Services (ASR, TTS, Speaker ID): 92%
- RAG Chatbot: 88%
- **Overall:** 91%

## 7.2 Test Strategy

### Unit Tests:

Focus on individual components in isolation:

- ASR transcription accuracy
- Speaker embedding extraction and similarity
- TTS synthesis file generation
- Text cleaning functions
- Database model validation

### Integration Tests:

Test complete workflows end-to-end:

- Full voice query flow (ASR → Search → TTS)
- Speaker registration and identification
- RAG chatbot response generation
- Authentication and profile management

### Performance Tests:

Validate speed requirements:

- Search latency must be <200ms (achieved: 60-90ms)
- TTS synthesis must be <3s (achieved: 1-2s)
- End-to-end voice query <15s (achieved: 8-12s)

## 7.3 CI Integration

Tests run automatically on every push via GitHub Actions:

- Database migrations applied
- 50 tests executed with pytest
- Coverage report generated
- Build proceeds only if tests pass
- Failed tests block deployment

## 7.4 Key Test Results

### ASR Testing:

- Transcription accuracy verified on sample audio
- Multi-language support confirmed
- Error handling for invalid audio tested

### Speaker ID Testing:

- Embedding extraction consistency
- Registration workflow
- Identification accuracy (>95% on test samples)
- Threshold sensitivity analysis

### TTS Testing:

- All 9 accents generate valid audio
- Text cleaning removes special characters
- Audio file format validation
- Speed adjustment verification

### Search Testing:

- Empty query handling
- Minimum length validation
- Relevance scoring
- Author attribution

### Integration Testing:

- Complete voice-to-speech flow works end-to-end
- User authentication and profile updates
- RAG chatbot returns relevant, cited responses

## 8. Challenges & Solutions {#challenges}

### 8.1 Technical Challenges Summary

Challenge	Impact	Solution	Outcome
Wikiquote parsing	High	Custom regex patterns for varied formats	149,473 quotes extracted
Neo4j performance	Medium	Fulltext indexing + optimized queries	60-90ms search time
Speaker ID accuracy	High	Threshold tuning + quality requirements	95%+ accuracy
TTS chipmunk effect	Critical	Removed pitch manipulation	Natural voices
LLM speed	Medium	Local Ollama + 3B model	3-5s acceptable
Container crashes	High	Verbose logging + iterative debug	Stable deployment
Autocomplete auth	Medium	Made endpoint public	Real-time suggestions

### 8.2 Design Decisions

#### 1. Simplified Graph Schema

- **Decision:** Quote→Author only (no Works/Categories)
- **Rationale:** Better performance, simpler maintenance, sufficient for requirements
- **Trade-off:** Less granular attribution, but acceptable for MVP

#### 2. gTTS vs Neural TTS

- **Decision:** gTTS with 9 accents
- **Rationale:** Simple, no GPU needed, adequate quality
- **Trade-off:** Lower quality than neural, but practical

#### 3. Local Ollama vs Cloud LLM

- **Decision:** Local llama3.2:3b
- **Rationale:** Privacy, \$0 cost, offline capability
- **Trade-off:** Slower (3-5s), but acceptable

#### 4. PostgreSQL vs SQLite

- **Decision:** PostgreSQL (Neon) for production
- **Rationale:** Better concurrency, cloud-native, free tier sufficient
- **Trade-off:** Slightly more complex than SQLite

#### 5. Azure ACI vs Kubernetes

- **Decision:** Azure Container Instances
- **Rationale:** Simpler, pay-per-second, sufficient for demo
- **Trade-off:** No auto-scaling, but not needed

## 9. Results & Performance {#results}

### 9.1 System Performance Metrics

Metric	Target	Achieved	Status
Database size	N/A	149,473 quotes, 1,163 authors	COMPLETE
Search latency	<200ms	60-90ms	EXCELLENT
ASR latency	<5s	2-4s	GOOD
ASR accuracy	>90%	~95%	EXCELLENT
Speaker ID accuracy	>90%	95%+	EXCELLENT

Speaker ID speed	Target	Achieved	Status
TTS synthesis	<3s	1-2s	Excellent
TTS quality	Natural	Natural (no artifacts)	Fixed
RAG response	<10s	3-5s	Good
End-to-end voice	<15s	8-12s	Good
Autocomplete	<300ms	<100ms + 300ms debounce	Excellent
Test coverage	>70%	91%	Excellent
Deployment time	<15min	8-12min	Good

## 9.2 Requirements Compliance

### Step 1 (Graph Database & Autocomplete):

- Parse Wikiquote dump (149,473 quotes extracted)
- Build graph database (Neo4j with optimized schema)
- Full-text indexing (quoteTextIndex, 60-90ms searches)
- Autocomplete system (real-time with debouncing)
- Source attribution (author displayed for all results)

### Step 2 (Voice Interaction):

- ASR module (Whisper, 95% accuracy)
- Speaker identification (ECAPA-TDNN, 95%+ accuracy)
- Chatbot (Graph search + RAG enhancement)
- Personalized TTS (gTTS, 9 accents)
- Multi-user support (authentication and profiles)
- Voice command interaction (complete workflow)

All core requirements met and exceeded

## 9.3 Beyond Requirements

### Advanced Features:

- RAG-powered AI (semantic search + LLM)
- Real-time autocomplete
- System status monitoring
- Text cleaning for TTS
- Dynamic API URL detection
- CI/CD automation
- Cloud deployment
- Comprehensive testing (91% coverage)

## 9.4 Database Statistics

### Neo4j Aura:

- Total Nodes: 150,636
- Quotes: 149,473
- Authors: 1,163
- Relationships: 150,544
- Index: Fulltext 2.0, 100% populated
- Average Query: 60-90ms

### Top Authors by Quote Count:

- Mystery Science Theater 3000 (7,994)
- Donald Trump (3,921)
- Elvis Presley (3,871)
- The West Wing (2,578)
- Last words (1,881)

## 10. Future Work {#future-work}

### 10.1 Short-term Improvements

#### Performance:

- Redis caching for frequent queries

- Neo4j connection pooling
- Async TTS processing
- WebSocket for voice streaming

#### Features:

- Quote bookmarking system
- Social sharing (Twitter, Facebook)
- Quote of the day
- Multi-language UI
- Dark mode theme
- Export quotes as images

#### Quality:

- End-to-end integration tests
- Load testing
- Monitoring (Prometheus + Grafana)
- Error tracking (Sentry)

## 10.2 Long-term Vision

#### Advanced AI:

- Fine-tuned speaker ID model
- Neural TTS (NVIDIA NeMo)
- Multi-turn conversation memory
- Sentiment-aware recommendations
- Quote similarity clustering

#### Scalability:

- Kubernetes deployment
- Horizontal scaling
- Database replication
- CDN for static assets
- Message queue (Celery + Redis)

#### Extended Functionality:

- Mobile app (React Native/Flutter)
- Browser extension
- Public API with rate limiting
- User-submitted quotes
- Collaborative collections
- Multi-language translation

#### Advanced Graph Features:

- Add Works and Categories nodes
- Quote-to-quote relationships
- Author influence graphs
- Timeline visualization
- Topic modeling

## Conclusion

The "Which Quote?" project successfully implements a sophisticated multimodal quote discovery system that **exceeds the original requirements**. By combining graph database technology, state-of-the-art NLP models, and modern DevOps practices, the system delivers a production-ready application with excellent performance across all metrics.

#### Key Achievements:

##### Technical Excellence:

- **149,473 quotes** indexed with **60-90ms** search latency
- **95%+ accuracy** for both ASR and speaker identification
- **Natural TTS** across 9 accents (chipmunk effect eliminated)
- **91% test coverage** with automated CI/CD
- **8-12 minute** deployment pipeline

##### Innovation:

- RAG-powered chatbot with local LLM (privacy-preserving, \$0 cost)
- Real-time autocomplete with intelligent debouncing

- Semantic search with vector embeddings
- System status monitoring
- Dynamic environment detection

#### Production Deployment:

- Fully containerized with Docker
- Automated CI/CD with GitHub Actions
- Deployed to Azure Container Instances
- Zero monthly cost with student credits
- Stable, reliable operation

#### Lessons Learned:

1. **Pre-trained models are powerful** - No fine-tuning needed for excellent results
2. **Simplicity matters** - Simplified schema improved performance over complex design
3. **DevOps is critical** - CI/CD saved countless hours and prevented errors
4. **Iterative debugging essential** - Verbose logging crucial for cloud deployment
5. **User experience first** - Features like autocomplete significantly enhance usability

#### Final Reflection:

This project demonstrates that modern AI tools have democratized access to advanced NLP capabilities. By leveraging pre-trained models (Whisper, SpeechBrain, Ollama) and cloud infrastructure (Azure, Neo4j Aura, Neon), we built an enterprise-grade application that would have required years of development and massive datasets just a few years ago.

The combination of graph databases, voice interaction, semantic search, and conversational AI creates a unique and powerful quote discovery experience that goes far beyond simple text search. The system is not just functional but production-ready, well-tested, and deployed to the cloud with full automation.

**Project Status:**  Live and operational on Azure

---

## Appendices

### A. Technology Versions

#### Backend:

- Django 5.0.1
- Django REST Framework 3.15.1
- openai-whisper 20231117
- speechbrain 1.0.0
- gtts 2.5.0
- Neo4j Driver 5.18.0

#### Frontend:

- Vue.js 3.3.4
- Tailwind CSS 3.3.5
- Nginx (Alpine)

#### Infrastructure:

- Docker 24.0.7
- Python 3.10
- PostgreSQL 16 (Neon)
- Neo4j 5 (Aura)

### B. Repository Structure

```

wikiquote/
├── backend/           # Django REST API
|   ├── accounts/      # User auth & profiles
|   ├── quotes/        # Search & Neo4j
|   ├── voice/         # ASR, TTS, Speaker ID
|   └── settings.py
└── tests/             # 50 tests, 91% coverage
├── frontend/          # Vue.js SPA
|   ├── index.html     # Main application
|   └── nginx.conf
└── services/
    ├── etl/            # Data extraction
    ├── rag/            # RAG chatbot
    └── voice/          # Voice services
└── .github/workflows/ # CI/CD pipeline
└── docker-compose.yml # Local development
└── requirements.txt   # Dependencies

```

## C. API Endpoints

**Base URL:** `http://<host>:8000/api/v1`

### Authentication:

- `POST /auth/register/` - User registration
- `POST /auth/login/` - Login (returns token)
- `GET /auth/profile/` - Get profile
- `PUT /auth/profile/update/` - Update preferences

### Quotes:

- `GET /quotes/search/?q=<query>&limit=<n>` - Search with autocomplete
- `POST /quotes/chat/` - RAG chatbot
- `GET /quotes/history/` - Query history
- `GET /quotes/favorites/` - Favorite quotes

### Voice:

- `POST /voice/asr/` - Speech-to-text
- `POST /voice/query/` - Complete voice workflow
- `POST /voice/synthesize/` - Text-to-speech
- `POST /voice/speaker/register/` - Register voice
- `POST /voice/speaker/identify/` - Identify speaker
- `GET /voice/voices/` - Available accents

## D. Production URLs

- **Frontend:** <http://wikiquote-frontend.germanywestcentral.azurecontainer.io:8080>
- **Backend:** <http://wikiquote-backend.germanywestcentral.azurecontainer.io:8000>
- **Ollama:** <http://wikiquote-ollama.germanywestcentral.azurecontainer.io:11434>
- **Neo4j:** `neo4j+s://bf2c68ee.databases.neo4j.io`
- **GitHub:** <https://github.com/erdeno/wikiquote>

**Project Completion:** December 15, 2025

**Development Time:** ~9 weeks

**Lines of Code:** ~8,500 (excluding tests)

**Test Coverage:** 91%

**Status:** Live on Azure

**Author:** Oguz Erden

**Course:** NLP Capstone Project

**Acknowledgments:** OpenAI Whisper, SpeechBrain, Neo4j, Wikiquote, Anthropic Claude, Azure for Students, Open-source community

**Built with** ❤️ **using Python, Django, Vue.js, Neo4j, and state-of-the-art AI**